

Getting Started with Elasticsearch

Apr.2022



Agenda



01

**What's
Elasticsearch?**



02

Use Cases



03

**Recent
Topics**

Takeaway



Takeaway


4

1. Elasticsearch (ES) is de facto standard for full-text search engines
2. There are some types of ES services, the main one is Elastic Stack
3. A similarity search system can be easily built in combination with Deep Learning



01

What's Elasticsearch?



- **Full-text search engine by Elastic co.**
 - De facto standard for full-text search engines
 - Search speed is faster than competitive services
- **Elasticsearch is adopted by many companies**
 - e.g.) Wikimedia, Facebook, Netflix, Quora, ...
- **Abundant complex search queries**
 - kNN logic can be add-on in ES Ver.7.3 later (Aug.2019)
 - Latest version is 8.1 (Apr.2022]



Why Elasticsearch?

7





- Very fast search speed
- Scalable
- Flexible search by simple query
- High affinity with ML/DL
- Faster development speed



Ref.) OSS Search Engines Comparison

8

Elasticsearch has a faster development speed

No.		1	2	3	4	5
OSS Name		Elasticsearch	Apache Solr	Splunk	NGT (Neighborhood Graph and Tree for Indexing)	SPTAG (Space Partition Tree And Graph)
Logo		 elasticsearch	 Solr	 splunk>	 NGT	--
Developer		Elastic	Apache Software Foundation	Splunk Inc.	Yahoo	Microsoft
1st Released		2010y	2004y	2003y	2016y	2019y
Github	Latest Version	Ver.8.1.2 (2022/04)	Ver.8.11.1 (2022/04)	?	v1.14.3 (2022/04)	1st Release
	#Contributors	1,729	275	--	11	26
	#Commits	63,686	35,803	--	139	126
Performance	Search Speed	◎	○	○	◎	◎
	Complex Search	○	△	○	△	△
	Scalable	◎	△	○	◎	◎
	Analytical Flexibility	◎	△	◎	--	--
	Affinity with ML	◎	△	△	◎	◎
Ranking	DB-Engines Ranking (As of Apr.2022)	7	20	13	--	--
	within Search Engine	1	3	2	--	--
	Popularity	◎	○	○	--	--



Service Type of Elasticsearch (3-types)

9

(a) Elastic Stack elastic stack

- OSS provided by Elastic co.
- Need a server to use it

(b) Elastic Cloud Elastic Cloud

- Full-managed cloud service provided by Elastic co.

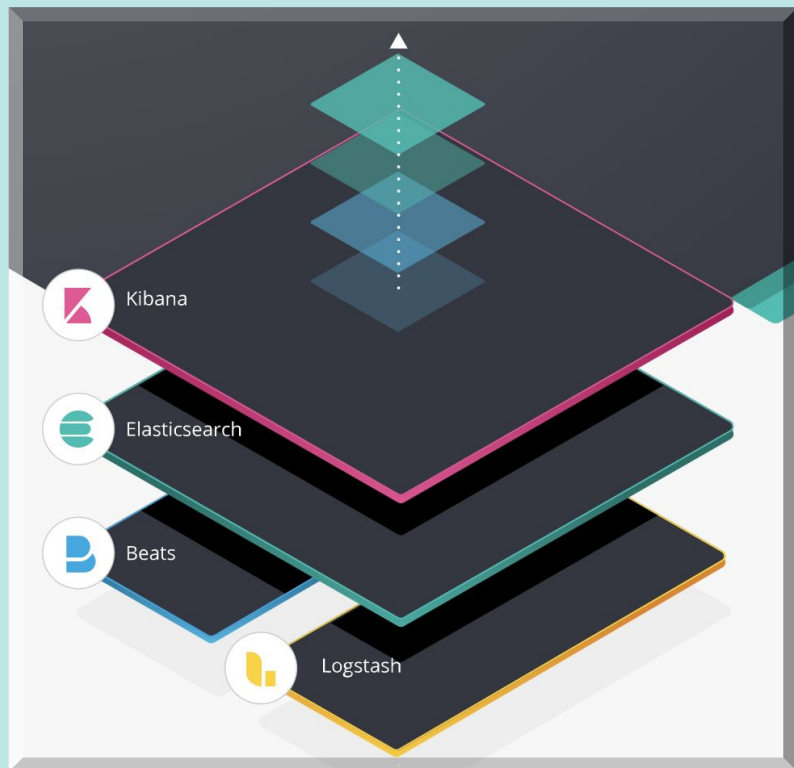
(c) Amazon OpenSearch Service OpenSearch

- Elasticsearch provided by AWS
- Old Service Name : “Amazon Elasticsearch Service” (~Sep.2021)

☞ Full-Managed (b) Elastic Cloud is recommended, if price is ignored!!

(a) Elastic Stack (also known as the ELK Stack)

10



- **Elasticsearch:**
Search Engine
- **Logstash:**
ETL Pipeline
- **Kibana:**
Searching/visualizing tool
- **Beats:**
Lightweight data shippers

(a) Elastic Stack: Elasticsearch

11

- Elasticsearch is a distributed, RESTful search and analytics engine capable of addressing a growing number of use cases
 - **“Ask your data questions of all kinds”**



The heart of the free and open Elastic Stack

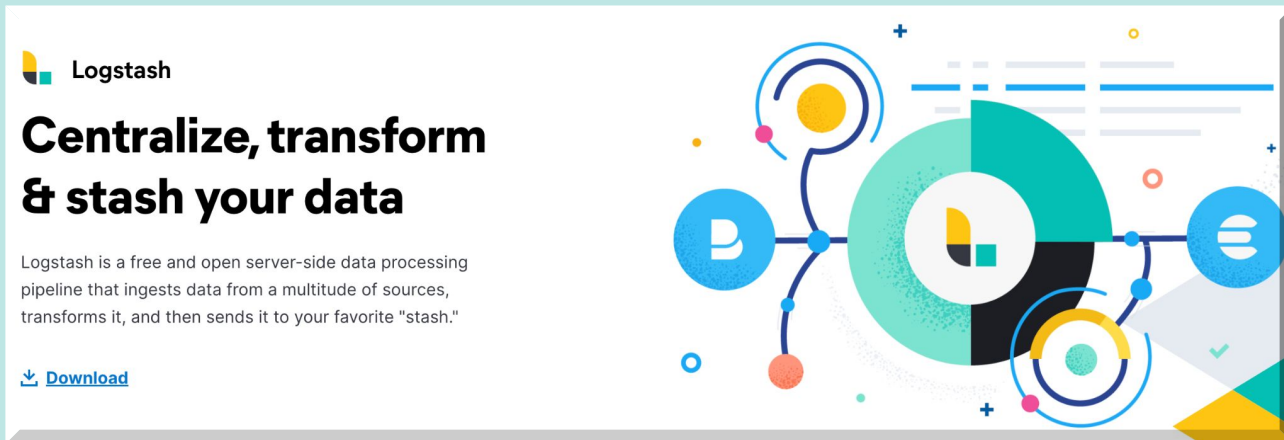
Elasticsearch is a distributed, RESTful search and analytics engine capable of addressing a growing number of use cases. As the heart of the Elastic Stack, it centrally stores your data for lightning fast search, fine-tuned relevancy, and powerful analytics that scale with ease.



(a) Elastic Stack: Logstash

12

- Logstash is a free and open server-side data processing pipeline that ingests data from a multitude of sources, transforms it, and then sends it to your favorite "stash."
 - **"Inputs, filters & outputs"**



The banner features the Logstash logo (a yellow and blue square) in the top left. Below it, the text "Centralize, transform & stash your data" is displayed in a bold, black font. Underneath this, a smaller line of text reads: "Logstash is a free and open server-side data processing pipeline that ingests data from a multitude of sources, transforms it, and then sends it to your favorite 'stash.'" At the bottom left, there is a blue "Download" link with a downward arrow icon. The right side of the banner is decorated with a colorful, abstract graphic consisting of a central teal circle, various blue and yellow circles, and connecting lines, suggesting a data pipeline or network.

Logstash

**Centralize, transform
& stash your data**

Logstash is a free and open server-side data processing pipeline that ingests data from a multitude of sources, transforms it, and then sends it to your favorite "stash."


[Download](#)



(a) Elastic Stack: Kibana

13

- Kibana is a free and open user interface that lets you visualize your Elasticsearch data and navigate the Elastic Stack
 - **“A picture's worth a thousand log lines”**

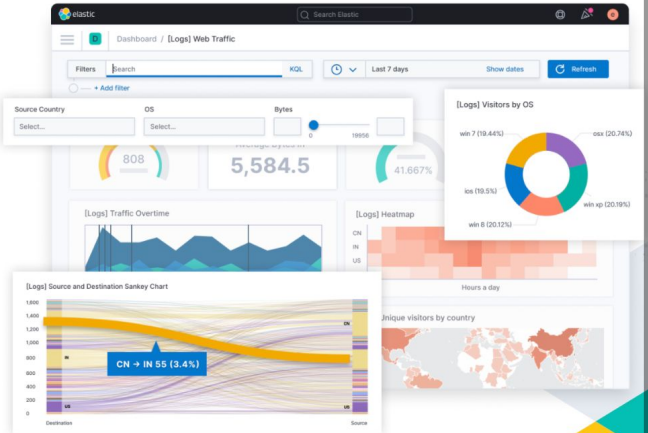
 **Kibana**

Your window into the Elastic Stack

Kibana is a free and open user interface that lets you visualize your Elasticsearch data and navigate the Elastic Stack. Do anything from tracking query load to understanding the way requests flow through your apps.

[Start free trial](#)

[Download Kibana](#)



Basic Terms in Elasticsearch

14

- Note that basic terms are different from RDB
 - The following terms correspond

RDB	Elasticsearch
database	index
table	mapping type
column	field
record	document



Exact match search and fuzzy search can be mixed

```
body = {  
  "from": skipN,  
  "size": topN,  
  "query": {  
    "bool": {  
      "filter": [  
        { 'term' : { 'xxx.keyword' : target_xxx } },  
        { 'term' : { 'yyy.keyword' : target_yyy } },  
      ],  
      "should" : [  
        { 'match' : { zzz : target_zzz } }  
      ]  
    }  
  },  
  "_source": {"includes": output_columns},  
  "sort" : [{"_score": "desc"}]  
}
```

Exact match
search



Fuzzy search



02

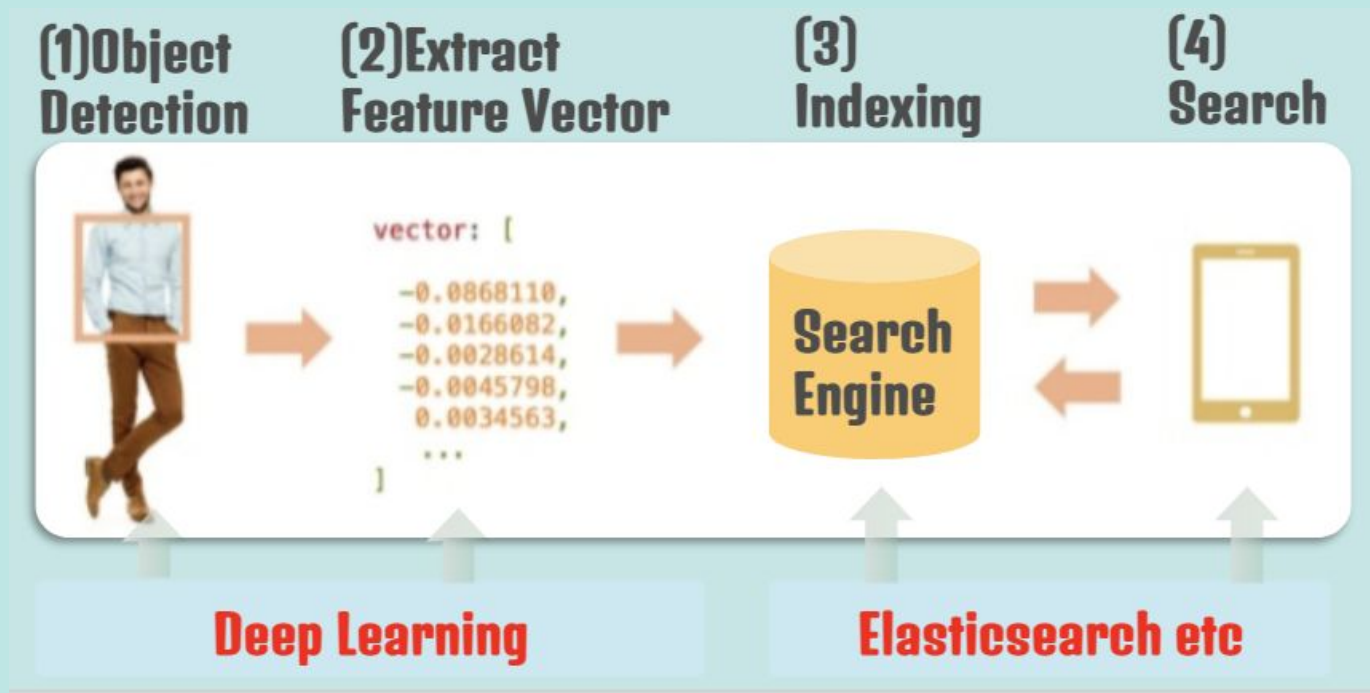
Use cases



Use-case(1): Image Similarity Search


17

Easily create similar image searches with DL + ES



Ex.) Company Clustering

Similarity search based on the tokenized texts of company profiles



Use cosine similarity between “embedding vectors”

```
body = {  
  'query': {  
    'script_score': {  
      'query': {  
        'bool': {  
          'filter': [  
            {'term': { 'xxx.keyword' : target_xxx }},  
            {'term': { 'yyy.keyword' : target_yyy }}, ],  
          'script': {  
            'source': ('cosineSimilarity(params.query_vector, doc["image_vector"]) + 1.0')*0.5,  
            'params': {'query_vector': query_vector}  
          }  
        }  
      }  
    },  
    '_source': {'includes': output_columns },  
    'sort' : [{'_score': 'desc'}], 'from': skipN, 'size': topN  
  }  
}
```

**exact match
search**

vector search

Normalize: 0.0~1.0
(Cosine: -1.0~1.0)

Embedding Vector (image/text \rightarrow numerical vector)

20

Word Embedding Vector refers to
“assign each word to a unique vector”

Embedding example

[apple, orange, banana]

"apple" = [1, 0, 0]

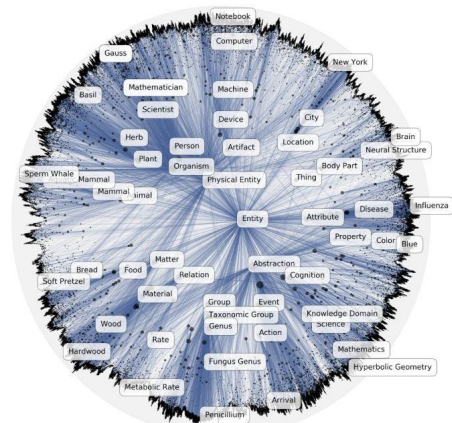
"banana" = [0, 0, 1]

All words are represented by the same dim.
So, it can be easily input as valuables to ML

Word embedding to some space

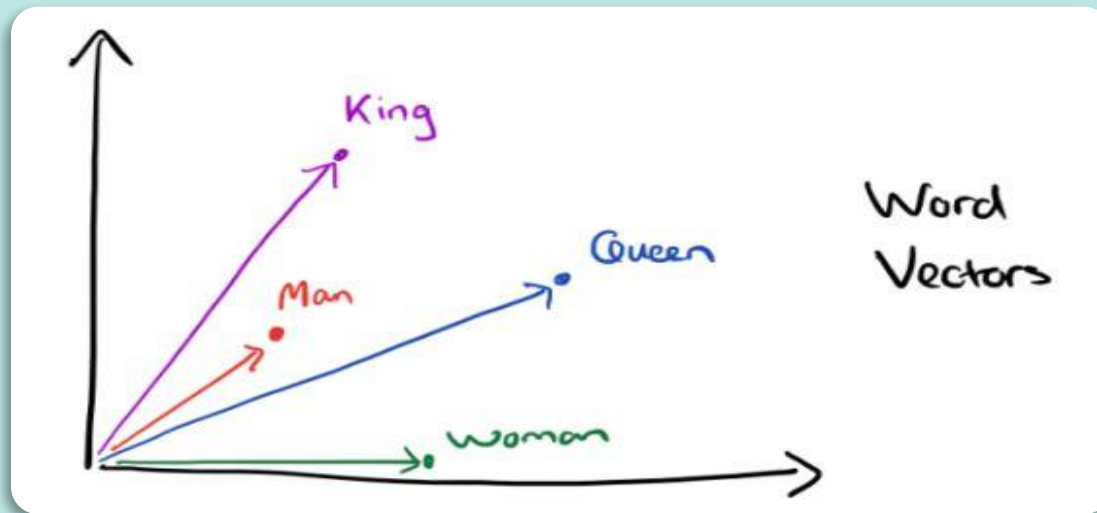
Poincaré Embeddings for Learning Hierarchical Representations

PyTorch implementation of Poincaré Embeddings for Learning Hierarchical Representations



[Facebook]: Poincaré Embeddings

King - Man + Woman = Queen!!



03

Recent Topics



Topic-01

Elastic vs AWS



Topic-01: Elastic vs AWS (1/4)

24




[\[Meet a CEO: Shay Banon from Elastic\]](#)

Topic-01: Elastic vs AWS (2/4)

25

Shay Banon: CTO(2012) -> CEO(2017) -> **CTO(2022)**




[Shay Banon to (re)assume the role of CTO,
Ash Kulkarni promoted to CEO (Jan.2022)]

Ref.) Elastic's Quarterly Revenue


26

- Elastic's quarterly revenue is growing consistently
- FY2021 Q1-Q4 sales totaled \$ 608.5M, YoY + 42% growth



■ Subscription ■ Services

Sums may not add to totals due to rounding.



[\[Investor Presentation and Company Overview\]](#)



Topic-01: Elastic vs AWS (3/4)

27

20 JANUARY 2021

NEWS

EN

ES

PT

KR

JP

DE

FR

CN

Amazon: NOT OK - why we had to change Elastic licensing

By [Shay Banon](#)

“

We have seen that this trademark issue drives confusion with users thinking Amazon Elasticsearch Service is actually a service provided jointly with Elastic, with our blessing and collaboration. This is just not true. NOT OK.

[\[Doubling down on open, Part II\]](#)

- Elastic changed to the dual license (Jan.2021)
 - **Server Side Public License (SSPL)**
 - **Elastic License** (restrict commercial services)
- The purpose is to prevent AWS from offering Elasticsearch and Kibana as managed services on their own





Topic-02 New Features of Ver.8.0

The major update in Feb. 2022 (7.x → 8.x)

- **Approximate k-nearest neighbor (A-kNN) searches**
 - So far, only exact kNN, which is hard to scale
 - Large-scale vector-based searches can be performed faster
- **Native support for NLP models**
 - Starting with Elastic Stack 8.0, NLP using external models published on HuggingFace has become very easy to run



Approximate k-nearest neighbor (A-kNN) searches

- So far, only exact kNN, which is hard to scale
- Large-scale vector-based searches can be performed faster

Common use cases for kNN:


- Relevance ranking based on NLP algorithms
- Product recommendations
- Similarity search for images or videos

[\[k-nearest neighbor \(kNN\) search\]](#)



Power of NLP to move analysis to unexplored territory

7.x



8.x



[Elastic 8.0: A new era of speed, scale, relevance, and simplicity]

We can easily try the pretrained NLP tasks

(1) Select a trained model (HuggingFace NLP tasks)

(2) Import the model and its tokenizer vocabulary

```
eland_import_hub_model  
--url <clusterUrl> \  
--hub-model-id elastic/[model identifier in the HuggingFace] \  
--task-type [NLP_task]
```

```
=== NLP tasks ===  
· fill_mask  
· ner(NER)  
· text_classification  
· text_embedding  
· zero_shot_classification
```

(3) Deploy the model in your cluster

(4) Inference (Try it out!!)

```
POST /_ml/trained_models/[NLP_task]/deployment/_infer  
{  
  "docs":{ "text_field" : "Sasha bought 300 shares of Acme Corp in 2022." }  
}
```



Bonus Slide



Is the trigger surprisingly simple?

35

Shay Banon created Elasticsearch while trying to index recipes for his wife, who was attending cooking school



[Elastic CEO Shay Banon stepping down to become CTO,
with Ash Kulkarni as new CEO]

End of Documents

